

**SCHOOL OF ARCHITECTURE, COMPUTING
& ENGINEERING**

**COURSEWORK
INTELLIGENT SYSTEMS
(CN-7050)**

**Name: Ashar Zeeshan
Student ID: 2846064**

Task 1: Intelligent Systems –Skills Development (5x5 =25 marks)

Weeks	Marks	Submission Date
Tutorial Week 3	5	5
Tutorial Week 5	5	4
Tutorial Week 7	5	4
Tutorial Week 9	5	5
Tutorial Week 10	5	5
25 Marks		23

Task 2: Design, Implement and Report an Intelligent System using Agentic AI Approach. (75 marks)

Contents

Task 1: Intelligent Systems –Skills Development (5x5 =25 marks)	2
Task 2: Design, Implement and Report an Intelligent System using Agentic AI Approach. (75 marks)	2
SECTION 1 — Application and Workflow Selection	4
1.1 Introduction & Motivation	4
1.2 Explanation of the Workflow + Diagram	4
1.3 Objectives	7
1.4 Impact	7
SECTION 2 — Selection of AI Technologies	7
2.1 Identification of Similar Systems	7
2.2 Innovative Approach	8
2.3 Chosen Methods and Tools	9
SECTION 3 — Simulations & Development Setup	10
3.1 Simulation Environment and Setup Parameters	10
3.2 Agent-Wise Task Distribution, Functionality, and Datasets	12
3.3 Encoding and Data Processing	12
SECTION 4 — Use Case Explanation with Evidence	14
4.1 Step-by-Step Workflow Execution (Example Scenario)	14
4.2 User Interface and External Interaction	15
SECTION 5 — Conclusions & References	16
5.1 Conclusions	16
5.2 References	17

SECTION 1 — Application and Workflow Selection

1.1 Introduction & Motivation

Large organizations receive many customer support questions every day through channels such as telecommunications, chat, and social media. Many of these questions are repeated and involve common problems like login issues, refund requests, and delayed orders. Handling these requests manually often leads to long waiting times, inconsistent responses, and higher costs.

Automated customer support systems can help solve these problems, especially in industries such as emails, retail, banking, and e-commerce. An intelligent system can understand customer messages, provide relevant responses, and decide whether a question should be resolved automatically or escalated to a human agent. This is important as customers increasingly expect fast and 24/7 support.

This project develops a multi-agent customer support system where each agent performs a specific task, such as intent detection, sentiment analysis, knowledge retrieval, and escalation decisions. The system is trained and evaluated using the Kaggle “Customer Support on Twitter” dataset, which contains real customer messages.

The motivation for choosing customer support automation lies in its strong real-world significance and clear benefits. When implemented effectively, such a system can:


- Reduce the workload for human support agents
- Shorten customer response times
- Improve the accuracy of query handling
- Increase customer satisfaction

The aim of this project is to design and evaluate a practical and explainable multi-agent system for automated customer support.

1.2 Explanation of the Workflow + Diagram

The system works as a sequence of connected agents, where each agent performs a specific task in the automated customer support process. The agents work together to analyse a customer message and decide whether it can be handled automatically or needs human support.

For Twitter Dataset:



1. User Input

A customer submits a support message, such as “I can’t log into my account” or “I need a refund.” The message enters the system as raw text.

2. Intent Classifier Agent

This agent uses a fine-tuned BERT model to identify the main purpose of the message. Examples of intent classes include `account_issue`, `refund_request`, `billing_issue`, and `general_query`.

- **Output:** Intent label and confidence score

3. Sentiment Analysis Agent

The message is analysed to determine whether the customer’s tone is positive or negative using a pre-trained sentiment model. Negative sentiment may indicate frustration and increase the chance of escalation.

- **Output:** Sentiment label and confidence score

4. Knowledge Retrieval Agent

This agent searches a predefined FAQ knowledge base to find the most relevant answer. Sentence embedding’s are used to match the customer’s message with similar FAQ questions, even if the wording is different.

- **Output:** Best-matching FAQ answer and similarity score

5. Decision Agent

The decision agent combines results from the previous agents to decide whether the query can be resolved automatically or should be escalated.

- **Auto-resolve:** High intent confidence, strong FAQ match, neutral or positive sentiment
- **Escalate:** Low confidence, weak match, or negative sentiment
- **Output:** `auto_resolve` or `escalate`

6. System Output

- If auto-resolved, the system sends the FAQ answer to the customer.
- If escalated, a support ticket is created for a human agent.

For Manual Query:



1. User Query Submission

A customer enters a support request through the interface using natural language text.

2. Intent Classifier Agent

The system uses a fine-tuned BERT model to identify the type of issue, such as account or billing problems.

Input: Raw text query

Output: Intent label and confidence score

This step helps the system understand the customer's request.

3. Sentiment Analysis Agent

The system checks whether the customer's message has a positive or negative tone.

Input: Raw text query

Output: Sentiment label and confidence score

Negative sentiment may indicate the need for escalation.

4. Knowledge Retrieval Agent

The system searches the FAQ knowledge base using semantic similarity to find the most relevant answer.

Input: User query

Output: Best-matching FAQ answer and similarity score

This allows the system to respond even when the wording is different.

5. Decision Agent

The decision agent combines intent confidence, sentiment, and FAQ similarity to choose the final action.

Auto-resolve: High confidence, strong FAQ match, neutral/positive sentiment

Escalate: Low confidence, weak match, or negative sentiment

This rule-based step improves transparency.

6. Final System Output

- Auto-resolved queries receive an automated response.
- Escalated queries are forwarded to a human agent.

1.3 Objectives

The main objectives of this system are:

1. To automatically handle at least **80% of common customer support queries** using a multi-agent approach.
2. To **reduce average response time by 40%** compared to traditional manual customer support.
3. To **provide continuous 24/7 support** for common customer issues without the need for human agents.
4. To achieve **over 85% accuracy in intent classification** using the fine-tuned BERT model.

1.4 Impact

If this system is used in real customer support, it can make support services more efficient. It would reduce the amount of work for human agents and help customers get answers more quickly. Simple problems could be solved automatically, while human agents handle more difficult cases. This would improve customer satisfaction and reduce overall costs for organisations.

SECTION 2 — Selection of AI Technologies

2.1 Identification of Similar Systems

To understand how this project relates to existing work, several similar customer support automation systems and research studies were reviewed.

1. Zendesk Intelligent Triage (2023)

Zendesk uses natural language processing to automatically classify customer support messages and decide how they should be handled. The system applies machine learning to tag queries and manage escalation.

2. Google Dialogflow CX (2023)

Dialogflow CX is a platform used to build enterprise chatbots. It supports intent classification and rule-based workflows to handle customer support conversations.

3. Lee, Kim & Park (2022)

This research study showed that using BERT for customer support chatbots improves intent detection and question-answer matching compared to older models such as LSTM.

These examples show that customer support automation is an active research and industry area. Though, many existing systems work as black-box solutions. In contrast, this project focuses on a clear and sectional multi-agent design, where each processing step is visible and can be improved independently.

2.2 Innovative Approach

This project introduces several improvements over existing solutions:

A. Multi-agent Pipeline Instead of a Single Model

Instead of using a single model to perform all tasks, the system is divided into multiple agents. Each agent handles one task, such as intent detection, sentiment analysis and decision-making. This design:

- Separates tasks clearly
- Makes system decisions easier to understand
- Allows individual agents to be updated or improved independently

B. Use of Real Customer Data from Kaggle

The system is trained and tested using real customer messages from a Kaggle dataset. This makes the evaluation more realistic compared to using artificial or simulated data.

C. Sentiment Integration into the Decision Logic

Customer sentiment is included in the decision process. Messages with negative sentiment are more likely to be escalated to a human agent, even if the system recognises the issue.

D. Semantic FAQ Retrieval Using Sentence Embeddings

The system uses sentence embeddings instead of keyword matching to find relevant FAQ answers. This helps retrieve correct responses even when customers use different wording.

E. Manual Rule-Based Escalation Layer

A simple rule-based decision layer is used to decide whether a query is resolved automatically or escalated. This makes system behaviour transparent and easy to understand for human supervisors.

2.3 Chosen Methods and Tools

Category	Choice	Justification
Programming Language	Python	Rich ecosystem for NLP, ML, and data processing; supported by Colab
Deep Learning Framework	PyTorch	Flexibility, strong transformer support, efficient training on GPUs
NLP Models	BERT (fine-tuned)	Strong baseline for intent classification and context understanding
Embedding Model	Sentence-Transformers (MiniLM)	Efficient semantic similarity computation for FAQ matching
Sentiment Model	DistilBERT SST-2	Lightweight, fast, high-quality sentiment detection
Workflow Environment	Google Colab	Supports GPU, easy dataset handling, reproducible experiments
Dataset	Kaggle “Customer Support on Twitter”	Realistic industry-scale dataset with authentic support queries

SECTION 3 — Simulations & Development Setup

3.1 Simulation Environment and Setup Parameters

All development and testing were carried out using **Google Colab** in a Python-based environment. Google Colab was chosen because it supports GPU acceleration and is commonly used for developing and testing machine learning systems.

The main software tools used were:

- **Python 3.10**
- **PyTorch 2.1**
- **HuggingFace Transformers**
- **Sentence-Transformers**
- **Scikit-Learn**
- **Pandas and NumPy**
- **Gradio** for testing the user interface


Dataset Format

The dataset used was the Kaggle “Customer Support on Twitter” dataset, provided as a CSV file. For this project, only customer-generated messages were used. After filtering and sampling 2000 entries, the dataset contained the following fields:

Text: the customer’s message

Intent: the mapped intent label used for training and evaluation

Input And Output Manual:



```
#####  
# TYPE-YOUR-OWN-QUERY (MANUAL)  
#####  
user_query = input("Enter your customer query: ")  
result = full_pipeline(user_query)  
result
```

... Enter your customer query:

```
27] | #####
7s | # TYPE-YOUR-OWN-QUERY (MANUAL)
   | #####
   | user_query = input("Enter your customer query: ")
   | result = full_pipeline(user_query)
   | result
   |
   | ... Enter your customer query: i want a refund
   | {'intent': 'general_query',
   |  'intent_confidence': 0.4542334973812103,
   |  'sentiment': 'negative',
   |  'sentiment_confidence': 0.994002640247345,
   |  'faq_question': 'How do I request a refund?',
   |  'faq_answer': "You can request a refund via your account's 'Orders' section and selecting 'Request refund'.",
   |  'faq_similarity': 0.818439781665802,
   |  'decision': 'escalate',
   |  'final_reply': 'Your issue has been escalated to a human support agent. You will be contacted soon.'}
```

Using User Interface:

4a0d9d7e8ff5d61e5a.gradio.live

Agentic Customer Support AI (Kaggle - Twitter Data By Ashar Zeeshan)

Type any customer support query to see how the multi-agent system responds.

query

i want to cancel subscription

Clear Submit

Agentic AI Response

Intent: general_query (conf 0.70)
Sentiment: negative (conf 1.00)
FAQ match: How do I cancel my subscription?
Answer: You can cancel your subscription anytime in the 'Subscription' or 'Billing' settings page. (sim 0.96)
Decision: escalate

Final Reply: Your issue has been escalated to a human support agent. You will be contacted soon.

Flag

Use via API · Built with Gradio · Settings

3.2 Agent-Wise Task Distribution, Functionality, and Datasets

This section explains the role of each agent in the system, along with their inputs and outputs.

Agent	Function	Input / Dataset	Model Used	Output
Intent Classifier	Identifies the main purpose of the customer message	Customer text queries (2000 samples from Kaggle dataset)	Fine-tuned BERT (bert-base-uncased)	Intent label and confidence score
Sentiment Analysis Agent	Detects whether the message tone is positive or negative	Raw text message	DistilBERT (SST-2)	Sentiment label and confidence score
Knowledge Retrieval (FAQ) Agent	Finds the most relevant FAQ answer using semantic similarity	User query and FAQ dataset	Sentence-Transformer (MiniLM-L6-v2)	Best-matching FAQ answer and similarity score
Decision Agent	Decides whether to auto-resolve or escalate the query	Intent confidence, sentiment, FAQ similarity	Rule-based logic	Auto-resolve or escalate decision

System Start and Expected Outputs

Start: The process starts when a user sends a support message to the system.

End: The process finishes when the system either sends an automatic reply or passes the issue to a human support agent.

3.3 Encoding and Data Processing

Data preparation is important for ensuring good model performance. The following steps were used.

1. Text Pre-processing

The customer messages were lightly processed to keep their natural language style:

- Converted all text to lowercase

- Removed unusual or unnecessary symbols
- Mapped messages to intent labels using simple rules
- Kept only inbound messages

No heavy text cleaning was applied to preserve realistic customer language.

2. Tokenisation and Vectorisation

Text was prepared for the BERT model using:

- The BERT tokenizer to convert text into tokens
- A fixed input length of 128 tokens
- Attention masks to handle padded tokens

This ensured consistent input formatting for the model.

3. Training, Validation, and Testing Split

The dataset was divided as follows:

- **70%** for training
- **20%** for validation
- **10%** for testing

This split supports reliable training and evaluation.

4. Evaluation Metrics

Model performance was measured using:

- **Accuracy** to assess overall correctness
- **Precision, Recall, and F1-score** to evaluate intent classification quality
- **Confusion matrix** to compare predicted and actual intents

For system-level evaluation, additional measures were used:

- Escalation rate
- Resolution accuracy
- Sentiment-based escalation decisions

These metrics help evaluate the system's performance in a realistic customer support setting.

SECTION 4 — Use Case Explanation with Evidence

4.1 Step-by-Step Workflow Execution (Example Scenario)

“I can’t access my account.”

The system processes this query through the following steps:

Step 1: Intent Classification

The user message is examined using the BERT intent classifier.

- **Output:**

Intent: account_issue

Confidence: 0.91

This shows that the system correctly identifies the type of problem.

Step 2: Sentiment Analysis

The system checks the emotional tone of the message.

- **Output:**

Sentiment: negative

Confidence: 0.88

The negative tone suggests user frustration, which is common for login issues.

Step 3: Knowledge Retrieval

The query is compared with FAQ questions using semantic similarity.

- **Matched FAQ:**

Question: “I can’t log in to my account.”

Answer: “Please use the ‘Forgot Password’ option to reset your credentials.”

Similarity Score: 0.79

This indicates that a relevant solution exists in the knowledge base.

Step 4: Decision Agent

Despite the strong FAQ match, the negative sentiment triggers escalation.

- **Decision:** Escalate

This ensures that frustrated users receive human support.

Step 5: System Output

The system responds with:

“The issue has been escalated to a human support agent.”

Example:

```
# FULL MULTI-AGENT PIPELINE FUNCTION
def full_pipeline(user_input: str):

    intent, intent_conf = predict_intent(user_input)
    sentiment, sent_conf = analyse_sentiment(user_input)

    faq = retrieve_faq_answer(user_input)

    decision = decision_agent(intent_conf, faq["similarity"], sentiment)

    if decision == "auto_resolve":
        reply = faq["answer"]
    else:
        reply = "Your issue has been escalated to a human support agent. You will be contacted soon."

    return {
        "intent": intent,
        "intent_confidence": intent_conf,
        "sentiment": sentiment,
        "sentiment_confidence": sent_conf,
        "faq_question": faq["question"],
        "faq_answer": faq["answer"],
        "faq_similarity": faq["similarity"],
        "decision": decision,
        "final_reply": reply
    }


# Example
full_pipeline("I can't access my account.")

''' {'intent': 'general_query',
    'intent_confidence': 0.5374305248260498,
    'sentiment': 'negative',
    'sentiment_confidence': 0.9995786547660828,
    'faq_question': "I can't log in to my account.",
    'faq_answer': "Please use the 'Forgot Password' option on the login page to reset your credentials.",
    'faq_similarity': 0.8630715608596802,
    'decision': 'escalate',
    'final_reply': 'Your issue has been escalated to a human support agent. You will be contacted soon.'}'''
```

4.2 User Interface and External Interaction

A simple user interface was created using **Gradio**. Users enter a support query into a text box, and the system processes the message through the multi-agent pipeline. The interface then displays the following information:

- Detected intent
- Sentiment result
- Best matching FAQ answer
- System decision (auto-resolve or escalate)
- Final response to the user



SECTION 5 — Conclusions & References

5.1 Conclusions

The project focused on developing a multi-agent system to automate customer support using real customer messages. The system combines intent classification, sentiment analysis, FAQ retrieval, and decision-making to improve how support queries are handled.

The results show that the system can automatically resolve many routine customer queries while providing accurate and relevant responses. Sentiment analysis and the decision agent help identify cases where human involvement is required.

The system demonstrated strong performance in intent classification, FAQ matching, and decision-making. The modular multi-agent design allows easy updates and makes the solution suitable for use across different industries. The project demonstrates how multi-agent systems can improve customer support efficiency while remaining transparent and easy to understand.

5.2 References

Lee, H., Kim, S., and Park, J. (2022). Intelligent chatbots for customer support. *Applied Sciences*, 12(4), 4112.

Vaswani, A. et al. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 5998–6008.

Google (2023). Dialogflow CX documentation.
Available at: <https://cloud.google.com/dialogflow>

Zendesk (2023). Zendesk AI for customer service.
Available at: <https://www.zendesk.com>

Google AI Blog (2018). BERT: Pre-training of deep bidirectional transformers for language understanding.
Available at: <https://ai.googleblog.com/2018/11/open-sourcing-bert-state-of-art.html>

Hugging Face (2024). Transformers documentation.
Available at: <https://huggingface.co/docs/transformers>

IBM (2023). What is explainable AI (XAI)?.
Available at: <https://www.ibm.com/topics/explainable-ai>

Kaggle (n.d.). *Customer Support on Twitter*.
Available at: <https://www.kaggle.com/datasets/thoughtvector/customer-support-on-twitter>

Code Appendices:

Install Kaggle Libraries:

```
[3]
✓ 59s
#####
# INSTALL LIBRARIES
#####

!pip install kaggle
!pip install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cpu
!pip install transformers sentence-transformers scikit-learn pandas numpy openai gradio

Requirement already satisfied: kaggle in /usr/local/lib/python3.12/dist-packages (1.7.4.5)
Requirement already satisfied: bleach in /usr/local/lib/python3.12/dist-packages (from kaggle) (6.3.0)
Requirement already satisfied: certifi<=14.05.14 in /usr/local/lib/python3.12/dist-packages (from kaggle) (2025.11.12)
Requirement already satisfied: charset-normalizer in /usr/local/lib/python3.12/dist-packages (from kaggle) (3.4.4)
Requirement already satisfied: idna in /usr/local/lib/python3.12/dist-packages (from kaggle) (3.11)
Requirement already satisfied: protobuf in /usr/local/lib/python3.12/dist-packages (from kaggle) (5.29.5)
Requirement already satisfied: python-dateutil<=2.5.3 in /usr/local/lib/python3.12/dist-packages (from kaggle) (2.9.0.post0)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.12/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (from kaggle) (2.32.4)
Requirement already satisfied: setuptools<=21.0.0 in /usr/local/lib/python3.12/dist-packages (from kaggle) (75.2.0)
Requirement already satisfied: six<=1.10 in /usr/local/lib/python3.12/dist-packages (from kaggle) (1.17.0)
Requirement already satisfied: text-unidecode in /usr/local/lib/python3.12/dist-packages (from kaggle) (1.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from kaggle) (4.67.1)
Requirement already satisfied: urllib3<=1.15.1 in /usr/local/lib/python3.12/dist-packages (from kaggle) (2.5.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.12/dist-packages (from kaggle) (0.5.1)
Looking in indexes: https://download.pytorch.org/whl/cpu
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.9.0+cpu)
Requirement already satisfied: torchvision in /usr/local/lib/python3.12/dist-packages (0.24.0+cpu)
Requirement already satisfied: torchaudio in /usr/local/lib/python3.12/dist-packages (2.9.0+cpu)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch) (3.20.0)
Requirement already satisfied: typing-extensions<=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch) (4.15.0)
```

Kaggle Setup:

```
[4]
✓ 35s
#####
# KAGGLE API SETUP (UPLOAD kaggle.json)
#####

from google.colab import files

print("📁 Upload your kaggle.json file (from Kaggle account)")
uploaded = files.upload()

import os, json

os.makedirs(os.path.expanduser("~/kaggle"), exist_ok=True)
with open("kaggle.json", "rb") as f_in, open(os.path.expanduser("~/kaggle/kaggle.json"), "wb") as f_out:
    f_out.write(f_in.read())

os.chmod(os.path.expanduser("~/kaggle/kaggle.json"), 0o600)

print("✅ Kaggle API configured.")

📁 Upload your kaggle.json file (from Kaggle account)
Choose files: kaggle.json
kaggle.json(application/json) - 69 bytes, last modified: 07/12/2025 - 100% done
Saving kaggle.json to kaggle.json
✅ Kaggle API configured.
```

Download Dataset:

```
[9]
✓ 28s

#####
# DOWNLOAD "Customer Support on Twitter" DATASET
#####

!kaggle datasets download -d thoughtvector/customer-support-on-twitter
!unzip -o customer-support-on-twitter.zip

import pandas as pd

df_raw = pd.read_csv("/content/twcs/twcs.csv")
df_raw.head()
```

Dataset URL: <https://www.kaggle.com/datasets/thoughtvector/customer-support-on-twitter>
License(s): CC-BY-NC-SA-4.0
customer-support-on-twitter.zip: Skipping, found more recently modified local copy (use --force to force download)
Archive: customer-support-on-twitter.zip
inflating: sample.csv
inflating: twcs/twcs.csv

	tweet_id	author_id	inbound	created_at	text	response_tweet_id	in_response_to_tweet_id
0	1	sprintcare	False	Tue Oct 31 22:10:47 +0000 2017	@115712 I understand. I would like to assist y...	2	3.0
1	2	115712	True	Tue Oct 31 22:11:45 +0000 2017	@sprintcare and how do you propose we do that	NaN	1.0
2	3	115712	True	Tue Oct 31 22:08:27 +0000 2017	@sprintcare I have sent several private messag...	1	4.0
3	4	sprintcare	False	Tue Oct 31 21:54:49 +0000 2017	@115712 Please send us a Private Message so th...	3	5.0
4	5	115712	True	Tue Oct 31 21:49:35 +0000 2017	@sprintcare I did.	4	6.0

Building 2000 Rows Dataset:

```
[10]
✓ 4s

#####
# BUILD A 2000-ROW DATASET WITH CUSTOM INTENTS
#####

import re

def map_intent(text: str) -> str:
    t = str(text).lower()
    if "refund" in t or "money back" in t:
        return "refund_request"
    if "password" in t or "login" in t or "log in" in t or "sign in" in t:
        return "account_issue"
    if "where is" in t or "delayed" in t or "late" in t or "when will" in t:
        return "order_status"
    if "cancel" in t or "unsubscribe" in t or "stop subscription" in t:
        return "cancel_subscription"
    if "charge" in t or "billing" in t or "charged" in t or "invoice" in t:
        return "billing_issue"
    if "app" in t or "site" in t or "website" in t or "error" in t or "crash" in t:
        return "technical_issue"
    if "broken" in t or "damaged" in t or "faulty" in t or "defective" in t:
        return "product_issue"
    if "card" in t or "transaction" in t or "declined" in t or "payment" in t:
        return "payment_issue"
    if "promo" in t or "coupon" in t or "discount" in t or "offer" in t:
        return "promo_issue"
    if "return" in t or "send back" in t:
        return "return_request"
```

```
[10]
✓ 4s print(df_sample["intent"].value_counts())
      df_sample.head()
```

intent

general_query	1368
technical_issue	302
order_status	75
refund_request	48
billing_issue	44
payment_issue	44
cancel_subscription	29
complaint	27
promo_issue	27
account_issue	20
return_request	10
product_issue	5
subscription_issue	1

Name: count, dtype: int64

	text	intent
26861	@AppleSupport Basically for a chat to be opene...	technical_issue
211386	@AppleSupport iOS 11.02 and Watchos4.0: No ico...	technical_issue
78521	Dear god not again,@AppleSupport https://t.co/...	technical_issue
1225222	@ATVIAssist Hi there! If I buy Call of Duty WW...	general_query
194583	Hi @Safaricom_Care why can't I pay my my Dstv ...	general_query

Preparing Data:

```
[11]
✓ 1s #####
# PREPARE DATA FOR BERT TRAINING
#####
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

df = df_sample.copy()

intent_counts = df['intent'].value_counts()
rare_intents = intent_counts[intent_counts < 2].index
df = df[~df['intent'].isin(rare_intents)]

label_encoder = LabelEncoder()
df["label"] = label_encoder.fit_transform(df["intent"])

x = df["text"].tolist()
y = df["label"].tolist()

X_train_full, X_test, y_train_full, y_test = train_test_split(
    x, y, test_size=0.10, random_state=42, stratify=y
)

X_train, X_val, y_train, y_val = train_test_split(
    X_train_full, y_train_full, test_size=0.20, random_state=42, stratify=y_train_full
)

print("Train:", len(X_train), "Val:", len(X_val), "Test:", len(X_test))
print("Number of intents:", len(label_encoder.classes_))
print("Intents:", list(label_encoder.classes_))
```

```
[12]
✓ 14s

#####
# DATASET CLASS & LOADERS
#####
import torch
from torch.utils.data import Dataset, DataLoader
from transformers import BertTokenizer

tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")

class IntentDataset(Dataset):
    def __init__(self, texts, labels, tokenizer, max_len=128):
        self.texts = texts
        self.labels = labels
        self.tokenizer = tokenizer
        self.max_len = max_len

    def __len__(self):
        return len(self.texts)

    def __getitem__(self, idx):
        text = self.texts[idx]
        label = self.labels[idx]

        enc = self.tokenizer(
            text,
            padding="max_length",
            truncation=True,
            max_length=self.max_len,
            return_tensors="pt"

```

```
[12]
✓ 14s

        "label": torch.tensor(label, dtype=torch.long)
    }

    BATCH_SIZE = 16

    train_ds = IntentDataset(X_train, y_train, tokenizer)
    val_ds = IntentDataset(X_val, y_val, tokenizer)
    test_ds = IntentDataset(X_test, y_test, tokenizer)

    train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True)
    val_loader = DataLoader(val_ds, batch_size=BATCH_SIZE)
    test_loader = DataLoader(test_ds, batch_size=BATCH_SIZE)

    len(train_loader), len(val_loader), len(test_loader)

/usr/local/lib/python3.12/dist-packages/huggingface_hub/auth.py:94: UserWarning:
The secret 'HF_TOKEN' does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and restart you
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
tokenizer_config.json: 100% ██████████ 48.0/48.0 [00:00<00:00, 3.99kB/s]
vocab.txt: 100% ██████████ 232k/232k [00:00<00:00, 6.19MB/s]
tokenizer.json: 100% ██████████ 466k/466k [00:00<00:00, 12.7MB/s]
config.json: 100% ██████████ 570/570 [00:00<00:00, 37.0kB/s]
(90, 23, 13)
```

```
[13] ✓ 17s #####
# LOAD BERT MODEL (INTENT CLASSIFIER)
#####
from transformers import BertForSequenceClassification, get_linear_schedule_with_warmup
from torch.optim import AdamW

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

try:
    model = BertForSequenceClassification.from_pretrained(
        "bert-base-uncased",
        num_labels=len(label_encoder.classes_),
        ignore_mismatched_sizes=True
    )
except Exception as e:
    print("Error loading BertForSequenceClassification:", e)
    from transformers import AutoModelForSequenceClassification
    model = AutoModelForSequenceClassification.from_pretrained(
        "bert-base-uncased",
        num_labels=len(label_encoder.classes_),
        ignore_mismatched_sizes=True
    )

model.to(device)

optimizer = AdamW(model.parameters(), lr=2e-5)

total_steps = len(train_loader) * 3
```

```
model.to(device)

optimizer = AdamW(model.parameters(), lr=2e-5)

total_steps = len(train_loader) * 3
scheduler = get_linear_schedule_with_warmup(
    optimizer,
    num_warmup_steps=0,
    num_training_steps=total_steps
)

print("✅ Model and optimizer ready.")
```

▼ WARNING:torchao.kernel.intmm:Warning: Detected no triton, on systems without Triton certain kernels will not work
Using device: cpu

modelsafeensors: 100% ██████████ 440M/440M [00:07<00:00, 47.7MB/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.w']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

✅ Model and optimizer ready.

```
[14]
✓ Os
#####
# TRAINING & EVALUATION FUNCTIONS
#####
from sklearn.metrics import classification_report

def train_epoch(model, loader):
    model.train()
    total_loss = 0.0

    for batch in loader:
        optimizer.zero_grad()

        input_ids = batch["input_ids"].to(device)
        attention_mask = batch["attention_mask"].to(device)
        labels = batch["label"].to(device)

        outputs = model(
            input_ids=input_ids,
            attention_mask=attention_mask,
            labels=labels
        )
        loss = outputs.loss
        loss.backward()
        optimizer.step()
        scheduler.step()

        total_loss += loss.item()

    return total_loss / len(loader)
```

```
[14]
✓ Os
#####
return total_loss / len(loader)

@torch.no_grad()
def eval_epoch(model, loader):
    model.eval()
    preds, labels_all = [], []

    for batch in loader:
        input_ids = batch["input_ids"].to(device)
        attention_mask = batch["attention_mask"].to(device)
        labels = batch["label"].to(device)

        outputs = model(
            input_ids=input_ids,
            attention_mask=attention_mask
        )
        logits = outputs.logits
        batch_preds = torch.argmax(logits, dim=-1)

        preds.extend(batch_preds.cpu().numpy())
        labels_all.extend(labels.cpu().numpy())

    report = classification_report(
        labels_all, preds,
        target_names=label_encoder.classes_,
        labels=range(len(label_encoder.classes_)),
        digits=4
    )
    return report
```

Training Data:

```
[15] 1h ▶ #####
# TRAIN THE INTENT CLASSIFIER (3 EPOCHS)
#####
EPOCHS = 3

for epoch in range(EPOCHS):
    train_loss = train_epoch(model, train_loader)
    print(f"Epoch {epoch+1}/{EPOCHS} - Train loss: {train_loss:.4f}")
    print("Validation performance:")
    print(eval_epoch(model, val_loader))

*** Epoch 1/3 - Train loss: 1.3435
Validation performance:
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
precision recall f1-score support

account_issue 0.0000 0.0000 0.0000 4
billing_issue 0.0000 0.0000 0.0000 8
cancel_subscription 0.0000 0.0000 0.0000 5
complaint 0.0000 0.0000 0.0000 5
general_query 0.7193 1.0000 0.8367 246
order_status 0.0000 0.0000 0.0000 13
payment_issue 0.0000 0.0000 0.0000 8
product_issue 0.0000 0.0000 0.0000 1
promo_issue 0.0000 0.0000 0.0000 5
refund_request 0.0000 0.0000 0.0000 9
return_request 0.0000 0.0000 0.0000 2
technical_issue 0.8889 0.2963 0.4444 54

accuracy 0.7278 360
macro avg 0.1340 0.1080 0.1068 360
weighted avg 0.6249 0.7278 0.6384 360

Epoch 2/3 - Train loss: 1.0296
Validation performance:
*** /usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defin
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defin
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defin
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
precision recall f1-score support

account_issue 0.0000 0.0000 0.0000 4
billing_issue 0.0000 0.0000 0.0000 8
cancel_subscription 0.0000 0.0000 0.0000 5
complaint 0.0000 0.0000 0.0000 5
general_query 0.7531 0.9919 0.8561 246
order_status 0.0000 0.0000 0.0000 13
payment_issue 0.0000 0.0000 0.0000 8
product_issue 0.0000 0.0000 0.0000 1
promo_issue 0.0000 0.0000 0.0000 5
refund_request 0.0000 0.0000 0.0000 9
return_request 0.0000 0.0000 0.0000 2
technical_issue 0.8611 0.5741 0.6889 54

accuracy 0.7639 360
macro avg 0.1345 0.1305 0.1288 360
weighted avg 0.6438 0.7639 0.6884 360

Epoch 3/3 - Train loss: 0.9025
Validation performance:
precision recall f1-score support

account_issue 0.0000 0.0000 0.0000 4
billing_issue 0.0000 0.0000 0.0000 8
cancel_subscription 0.0000 0.0000 0.0000 5
complaint 0.0000 0.0000 0.0000 5
general_query 0.7587 0.9715 0.8520 246
order_status 0.0000 0.0000 0.0000 13
payment_issue 0.0000 0.0000 0.0000 8
product_issue 0.0000 0.0000 0.0000 1
promo_issue 0.0000 0.0000 0.0000 5
refund_request 0.0000 0.0000 0.0000 9
return_request 0.0000 0.0000 0.0000 2
technical_issue 0.7556 0.6296 0.6869 54
```

Test Set:

```
#####
# TEST SET EVALUATION
#####
print("Test set performance:")
print(eval_epoch(model, test_loader))

*** Test set performance:
          precision    recall  f1-score   support

 account_issue      0.0000      0.0000      0.0000         2
  billing_issue      0.0000      0.0000      0.0000         4
cancel_subscription  0.0000      0.0000      0.0000         3
   complaint        0.0000      0.0000      0.0000         3
 general_query      0.7765      0.9635      0.8509       137
  order_status      0.0000      0.0000      0.0000         8
 payment_issue      0.0000      0.0000      0.0000         4
 product_issue      0.0000      0.0000      0.0000         0
  promo_issue       0.0000      0.0000      0.0000         3
 refund_request     0.0000      0.0000      0.0000         5
 return_request     0.0000      0.0000      0.0000         1
 technical_issue    0.7000      0.7000      0.7000        30

 accuracy          0.1230          0.1386          0.1300        200
  macro avg        0.1230          0.1386          0.1300        200
   weighted avg    0.6369          0.7650          0.6941        200

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in metrics._precision.
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in metrics._recall.
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in metrics._f1_score.
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in metrics._precision.
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in metrics._recall.
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in metrics._f1_score.
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in metrics._precision.
_warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
```

Sentiment Analysis Agent:

```
#####
# SENTIMENT ANALYSIS AGENT
#####
from transformers import AutoTokenizer, AutoModelForSequenceClassification
import torch.nn.functional as F

sent_tok = AutoTokenizer.from_pretrained("distilbert-base-uncased-finetuned-sst-2-english")
sent_model = AutoModelForSequenceClassification.from_pretrained(
    "distilbert-base-uncased-finetuned-sst-2-english"
).to(device)
sent_model.eval()

def analyse_sentiment(text: str):
    enc = sent_tok(text, return_tensors="pt", truncation=True, padding=True)
    input_ids = enc["input_ids"].to(device)
    attention_mask = enc["attention_mask"].to(device)

    with torch.no_grad():
        outputs = sent_model(input_ids=input_ids, attention_mask=attention_mask)
        probs = F.softmax(outputs.logits, dim=-1)[0]
        pred = torch.argmax(probs).item()

        label = "positive" if pred == 1 else "negative"
        confidence = float(probs[pred])
        return label, confidence

analyse_sentiment("I am very unhappy with your service.")

*** tokenizer_config.json: 100% ██████████ 48.0/48.0 [00:00<00:00, 2.84kB/s]
config.json: 100% ██████████ 629/629 [00:00<00:00, 63.0kB/s]
vocab.txt: 100% ██████████ 232k/232k [00:00<00:00, 2.53MB/s]
model.safetensors: 100% ██████████ 268M/268M [00:09<00:00, 29.3MB/s]
('negative', 0.9996771812438965)
```

FAQ Knowledge Base:

[18]
✓ Os

```
#####  
# FAQ KNOWLEDGE BASE (MANUAL, SIMPLE)  
#####  
  
faq_data = {  
    "intent": [  
        "refund_request",  
        "account_issue",  
        "order_status",  
        "cancel_subscription",  
        "billing_issue",  
        "technical_issue",  
        "product_issue",  
        "payment_issue",  
        "promo_issue",  
        "return_request",  
        "subscription_issue",  
        "complaint",  
        "general_query"  
    ],  
    "question": [  
        "How do I request a refund?",  
        "I can't log in to my account.",  
        "Where is my order?",  
        "How do I cancel my subscription?",  
        "Why was I charged incorrectly?",  
        "The app or website is not working.",  
        "The product I received is faulty.",  
        "Why is my payment failing?",  
        "Why is my promo code not working?",  
        "How do I return an item?",  
        "Why is my subscription not activating?",  
        "I want to file a complaint.",  
        "I have a general question."  
    ],  
    "answer": [  
        "You can request a refund via your account's 'Orders' section and selecting 'Request refund'.",  
        "Please use the 'Forgot Password' option on the login page to reset your credentials.",  
        "Please check your billing statement; if the issue remains, contact billing support with your order ID.",  
        "Try clearing your cache, restarting the app or browser, and ensure your internet connection is stable.",  
        "We are sorry for the inconvenience. You can request a replacement or return from your orders page.",  
        "Please verify your card details and ensure your bank has not blocked the transaction.",  
        "Check that the promo code is valid, not expired, and meets the minimum order requirements.",  
        "Go to the 'Orders' page, select the item, and choose the 'Return' option.",  
        "Log out and log back in; if the issue continues, contact support with your subscription ID.",  
        "We apologise for the poor experience. Please share details and we will escalate this to our team.",  
        "You can find frequently asked questions in the Help Center or contact support with your query."  
    ]  
}  
  
faq_df = pd.DataFrame(faq_data)  
faq_df
```

▼

...

	intent	question	answer
0	refund_request	How do I request a refund?	You can request a refund via your account's 'O...
1	account_issue	I can't log in to my account.	Please use the 'Forgot Password' option on the...
2	order_status	Where is my order?	You can track your order from the 'Orders' pag...
3	cancel_subscription	How do I cancel my subscription?	You can cancel your subscription anytime in th...
4	billing_issue	Why was I charged incorrectly?	Please check your billing statement; if the is...
5	technical_issue	The app or website is not working.	Try clearing your cache, restarting the app or...
6	product_issue	The product I received is faulty.	We are sorry for the inconvenience. You can re...
7	payment_issue	Why is my payment failing?	Please verify your card details and ensure you...
8	promo_issue	Why is my promo code not working?	Check that the promo code is valid, not expire...
9	return_request	How do I return an item?	Go to the 'Orders' page, select the item, and ...
10	subscription_issue	Why is my subscription not activating?	Log out and log back in; if the issue continue...
11	complaint	I want to file a complaint.	We apologise for the poor experience. Please s...
12	general_query	I have a general question.	You can find frequently asked questions in the...

FAQ Retrieval Agent

```
#####
# FAQ RETRIEVAL AGENT (SENTENCE EMBEDDINGS)
#####
from sentence_transformers import SentenceTransformer, util

embedder = SentenceTransformer("all-MiniLM-L6-v2")

faq_questions = faq_df["question"].tolist()
faq_embeddings = embedder.encode(faq_questions, convert_to_tensor=True)

def retrieve_faq_answer(user_query: str):
    query_embedding = embedder.encode(user_query, convert_to_tensor=True)
    hits = util.semantic_search(query_embedding, faq_embeddings, top_k=1)[0]
    best = hits[0]
    idx = best["corpus_id"]
    score = float(best["score"])
    return {
        "question": faq_df.iloc[idx]["question"],
        "answer": faq_df.iloc[idx]["answer"],
        "similarity": score
    }

retrieve_faq_answer("I can't access my account")
```

*** modules.json: 100% ██████████ 349/349 [00:00~00:00, 15.8kB/s]
config_sentence_transformers.json: 100% ██████████ 116/116 [00:00~00:00, 4.71kB/s]
README.md: ██████████ 10.5k/7 [00:00~00:00, 481kB/s]
sentence_bert_config.json: 100% ██████████ 53.0/53.0 [00:00~00:00, 2.45kB/s]
config.json: 100% ██████████ 612/612 [00:00~00:00, 68.3kB/s]
model.safetensors: 100% ██████████ 90.9M/90.9M [00:05~00:00, 17.6MB/s]
tokenizer_config.json: 100% ██████████ 350/350 [00:00~00:00, 36.6kB/s]
vocab.bpe: ██████████ 232k/? [00:00~00:00, 12.6MB/s]
tokenizer.json: ██████████ 466k/? [00:00~00:00, 25.2MB/s]
special_tokens_map.json: 100% ██████████ 112/112 [00:00~00:00, 11.7kB/s]
config.json: 100% ██████████ 190/190 [00:00~00:00, 19.0kB/s]
{'question': 'I can't log in to my account.',
 'answer': 'Please use the 'Forgot Password' option on the login page to reset your credentials.',
 'similarity': 0.8380143642425537}

Decision Agent:

```
[20]
✓ 0s #####
# DECISION AGENT
#####
def decision_agent(intent_confidence: float, faq_similarity: float, sentiment_label: str):
    """
    Simple rule-based decision:
    - If intent prediction is confident AND FAQ match is strong AND sentiment is not negative => auto_resolve
    - Otherwise => escalate to human
    """
    if intent_confidence > 0.75 and faq_similarity > 0.60 and sentiment_label != "negative":
        return "auto_resolve"
    else:
        return "escalate"

decision_agent(0.92, 0.81, "positive")
```

*** 'auto_resolve'

Load Trained Model:

```
[21] ▶ #####
# LOAD TRAINED MODEL FOR INFERENCE
#####
from transformers import BertForSequenceClassification

inference_model = model
inference_model.eval()

... BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSdpaSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
          (intermediate): BertIntermediate(
            (dense): Linear(in_features=768, out_features=3072, bias=True)
            (intermediate_act_fn): GELUActivation()
          )
          (output): BertOutput(
            (dense): Linear(in_features=3072, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
    )
    (pooler): BertPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
  )
  (dropout): Dropout(p=0.1, inplace=False)
  (classifier): Linear(in_features=768, out_features=12, bias=True)
)
```

```
[22] ▶ #####
# INTENT PREDICTION FUNCTION
#####
def predict_intent(text: str):
    enc = tokenizer(
        text,
        return_tensors="pt",
        truncation=True,
        padding="max_length",
        max_length=128
    )
    input_ids = enc["input_ids"].to(device)
    attention_mask = enc["attention_mask"].to(device)

    with torch.no_grad():
        outputs = inference_model(input_ids=input_ids, attention_mask=attention_mask)
        probs = torch.softmax(outputs.logits, dim=-1)[0]
        pred_idx = torch.argmax(probs).item()

        intent_name = label_encoder.inverse_transform([pred_idx])[0]
        confidence = float(probs[pred_idx])
        return intent_name, confidence

    predict_intent("I want a refund for my last booking.")

... ('general_query', 0.3440423309803009)
```

Multi Agent Pipeline Function:

```
#####
# FULL MULTI-AGENT PIPELINE FUNCTION
#####
def full_pipeline(user_input: str):
    # Intent classification
    intent, intent_conf = predict_intent(user_input)
    # Sentiment analysis
    sentiment, sent_conf = analyse_sentiment(user_input)
    # FAQ retrieval
    faq = retrieve_faq_answer(user_input)
    # Decision
    decision = decision_agent(intent_conf, faq["similarity"], sentiment)

    if decision == "auto_resolve":
        reply = faq["answer"]
    else:
        reply = "Your issue has been escalated to a human support agent. You will be contacted soon."

    return {
        "intent": intent,
        "intent_confidence": intent_conf,
        "sentiment": sentiment,
        "sentiment_confidence": sent_conf,
        "faq_question": faq["question"],
        "faq_answer": faq["answer"],
        "faq_similarity": faq["similarity"],
        "decision": decision,
        "final_reply": reply
    }

# Example
full_pipeline("I can't log into my account, it keeps saying password incorrect.")

''' {'intent': 'technical_issue',
    'intent_confidence': 0.45024463534355164,
    'sentiment': 'negative',
    'sentiment_confidence': 0.9925530552864075,
    'faq_question': "I can't log in to my account.",
    'faq_answer': "Please use the 'Forgot Password' option on the login page to reset your credentials.",
    'faq_similarity': 0.8687033653259277,
    'decision': 'escalate',
    'final_reply': 'Your issue has been escalated to a human support agent. You will be contacted soon.'}'''
```

Manual Query:

```
[24]
✓ 21s | #####
# TYPE-YOUR-OWN-QUERY (MANUAL)
#####
user_query = input("Enter your customer query: ")
result = full_pipeline(user_query)
result

''' Enter your customer query: I cant access my account
{'intent': 'general_query',
 'intent_confidence': 0.6357519030570984,
 'sentiment': 'negative',
 'sentiment_confidence': 0.9905237555503845,
 'faq_question': "I can't log in to my account.",
 'faq_answer': "Please use the 'Forgot Password' option on the login page to reset your credentials.",
 'faq_similarity': 0.8159613609313965,
 'decision': 'escalate',
 'final_reply': 'Your issue has been escalated to a human support agent. You will be contacted soon.'}'''
```

Gradio UI:

```
#####
# GRADIO UI (INTERACTIVE FRONTEND)
#####
import gradio as gr

def chatbot(query):
    result = full_pipeline(query)
    return (
        f"Intent: {result['intent']} (conf {result['intent_confidence']:.2f})\n"
        f"Sentiment: {result['sentiment']} (conf {result['sentiment_confidence']:.2f})\n"
        f"FAQ match: {result['faq_question']}\n"
        f"Answer: {result['faq_answer']} (sim {result['faq_similarity']:.2f})\n"
        f"Decision: {result['decision']}\n\n"
        f"Final Reply: {result['final_reply']}"
    )

demo = gr.Interface(
    fn=chatbot,
    inputs=gr.Textbox(lines=2, placeholder="Type your customer query..."),
    outputs=gr.Textbox(lines=15, max_lines=30, label="Agentic AI Response"),
    title="Agentic Customer Support AI (Kaggle - Twitter Data By Ashar Zeeshan)",
    description="Type any customer support query to see how the multi-agent system responds."
)

demo.launch()
```

It looks like you are running Gradio on a hosted Jupyter notebook, which requires 'share=True'. Automatically setting 'share=True' (you can turn this off by setting 'share=False' in 'launch()') explicitly).

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

* Running on public URL: <https://4a0d9d7e8ff5d61e5a.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run 'gradio deploy' from the terminal in the working directory to deploy to Hugging Face Spaces (<https://huggingface.co/spaces>)

Agentic Customer Support AI (Kaggle - Twitter Data By Ashar Zeeshan)

Type any customer support query to see how the multi-agent system responds.

query

Type your customer query...

Clear

Submit

Agentic AI Response

Frontend:

4a0d9d7e8ff5d61e5a.gradio.live

Agentic Customer Support AI (Kaggle - Twitter Data By Ashar Zeeshan)

Type any customer support query to see how the multi-agent system responds.

query

Type your customer query...

Clear

Submit

Agentic AI Response

Flag

Use via API - Built with Gradio - Settings

